

Exercise 5 Report/Write-Up

Dhiraj Bennadi and Patrick Bales-Parks

April 2, 2022

ECEN 5623- Spring 2022

1 Problem 1

1.1 problem

Download seqgen.c and seqgen2x.c (or unzip them from the provided zip file for this exercise) and build them in Linux on the Altera DE1-SOC, Raspberry Pi or Jetson board and execute the code. Describe what it is doing and make sure you understand how to use it to structure an embedded system. Determine the worst case execution time (WCET) for each service by printing or logging timestamps between two points in your code or by use of a profiling tool. Determine D, T, and C for each service and create an RM schedule in Cheddar using your WCET estimates. Calculate the percentage CPU utilization for this system.

1.2 Answer

As was suggested in the comments of the seqgen files, cpu check were preformed on the Jetson Nano prior to initialization.

```
patrick@rtnano: ~/exercise5/e5
make: *** [seqgen.o] Error 1
patrick@rtnano:~/exercise5/e5$ make
gcc -O0 -g -lm -c seqgen.c
gcc -O0 -g -lm -o seqgen seqgen.o -lpthread -lrt
patrick@rtnano:~/exercise5/e5$ lscpu
Architecture:        aarch64
Byte Order:          Little Endian
CPU(s):              4
On-line CPU(s) list: 0-3
Thread(s) per core:  1
Core(s) per socket:  4
Socket(s):           1
Vendor ID:           ARM
Model:               1
Model name:          Cortex-A57
Stepping:            r1p1
CPU max MHz:         1479.0000
CPU min MHz:         102.0000
BogoMIPS:            38.40
L1d cache:           32K
L1i cache:           48K
L2 cache:            2048K
Flags:               fp asimd evtstrm aes pmull sha1 sha2 crc32
patrick@rtnano:~/exercise5/e5$
```

Figure 1: Jetson Nano lscpu

1.2.1 Explanation and Execution of seqgen.c

The seqgen.c file was written to demonstrate the use of sequences in a Real-Time system. The program is designed to call the sequencer function as the main thread with highest priority. The sequences are called as threads, inside the main thread, with priorities lower than the main sequencer thread. These priorities are determined by the RM policy. in the original state, the sequencer operates at 30Hz. The sequencer loops through the services, only activating them at the designated service frequency. These frequencies are: $s_1 = 3\text{Hz}$, $s_2 = 1\text{Hz}$, $s_3 = 0.5\text{Hz}$, $s_4 = 1\text{Hz}$, $s_5 = 0.5\text{Hz}$, $s_6 = 1\text{Hz}$ and $s_7 = 0.1\text{Hz}$. The sequences are also protected by individual semaphores. Execution of seqgen.c and the values used for analysis can be seen in the figures below.

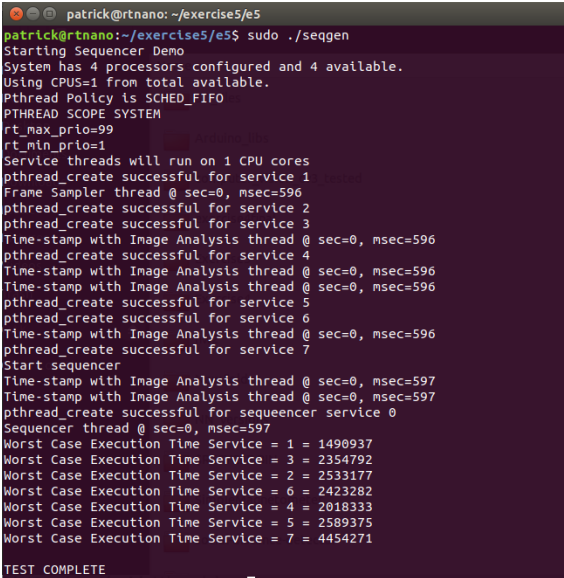


Figure 2: seqgen.c execution

Service	Frequency	Time Period T in ms	Deadline D, D=T	Priority	WCET	WCET in ms	Execution Time C
S1	3	333.3333333	333.3333333	1	1266302	1.266302	1.266302
S2	1	1000	1000	2	1941875	1.941875	1.941875
S4	1	1000	1000	2	2165520	2.16552	2.16552
S6	1	1000	1000	2	1752448	1.752448	1.752448
S3	0.5	2000	2000	3	1831250	1.83125	1.83125
S5	0.5	2000	2000	3	2475833	2.475833	2.475833
S7	0.1	10000	10000	4	1628490	1.62849	1.62849

Figure 3: seqgen.c values

1.2.2 Explanation and Execution of seqgen2.c

The seqgen2.c program demonstrates the same patterns as seqgen.c, but is adapted to operate at twice the frequency of 60Hz. Since the primary thread is the sequencer function, this is the thread changed from 30Hz to 60Hz. This change is accomplished by changing the "timespec delaytime" from 333333333 to 166666666. This changes the delay in nano seconds equivalent to 60Hz. Changing this parameter also means that the service frequencies had to be adjusted. The new service frequencies are: s1 = 30Hz, s2 = 10Hz, s3 = 5Hz, s4 = 10Hz, s5 = 5Hz, s6 = 10Hz and s7 = 1Hz. Execution of seqgen2.c and the values used for analysis can be seen in the figures below.

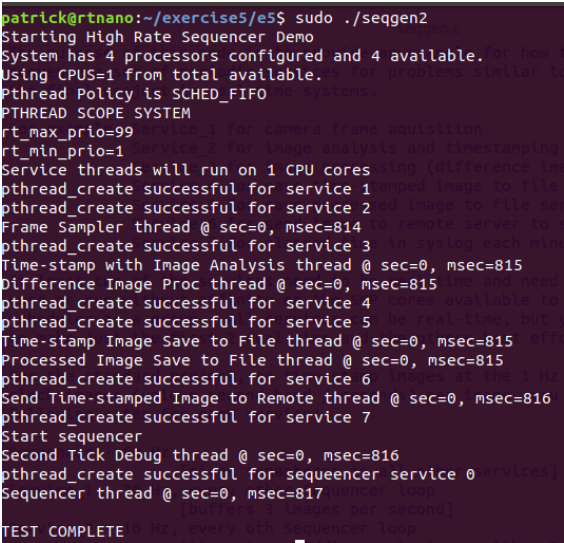


Figure 4: seqgen.c execution

Service	Frequency	Time Period T in ms	Deadline D, D=T	Priority	WCET	WCET in ms	Execution Time C
S1	30	33.33333333	33.33333333	1	1395938	1.395938	1.395938
S2	10	100	100	2	1943073	1.943073	1.943073
S4	10	100	100	2	2298021	2.298021	2.298021
S6	10	100	100	2	1749011	1.749011	1.749011
S3	5	200	200	3	1430416	1.430416	1.430416
S5	5	200	200	3	1684479	1.684479	1.684479
S7	1	1000	1000	4	2037500	2.0375	2.0375

Figure 5: seqgen.c values

1.2.3 Cheddar Schedule and CPU Utilization

Cheddar does not allow for the use of decimal values in its parameters. To compensate for this, the values used for execution time C (WCET) were rounded up to the next integer. With a higher execution time, the lower value of the true execution time is still guaranteed to succeed. The total CPU Utilization for each program can be seen in the figure below, followed by the Cheddar simulation task parameters and results.

seqgen.c cpu utilization				seqgen2.c cpu utilization			
Service	T	C	U	Service	T	C	U
S1	333.333	1.266	0.0038	S1	33.3333	1.396	0.04188
S2	1000	1.942	0.00194	S2	100	1.943	0.01943
S4	1000	2.165	0.00217	S4	100	2.298	0.02298
S6	1000	1.752	0.00175	S6	100	1.749	0.01749
S3	2000	1.831	0.00092	S3	200	1.43	0.00715
S5	2000	2.476	0.00124	S5	200	1.684	0.00842
S7	10000	1.629	0.00016	S7	1000	2.038	0.002038
Total Utilization =			1.1973%	Total Utilization =			11.9388%

Figure 6: Total CPU Utilization

The Cheddar simulation task parameters and results can be seen here.

Name	Task Type	Processor Name	Address Space	Capacity	Deadline	Start time	Prior
S1	Periodic	cpu1	Add0	1	3	0	1
S2	Periodic	cpu1	Add0	2	10	0	2
S3	Periodic	cpu1	Add0	2	20	0	3
S4	Periodic	cpu1	Add0	2	10	0	2
S5	Periodic	cpu1	Add0	2	20	0	3
S6	Periodic	cpu1	Add0	2	10	0	2
S7	Periodic	cpu1	Add0	2	100	0	4

Figure 7: Tasks for seqgen.c

Name	Task Type	Processor Name	Address Space	Capacity	Deadline	Start time	Prior
S1	Periodic	cpu1	Add0	2	33	0	1
S2	Periodic	cpu1	Add0	2	100	0	2
S3	Periodic	cpu1	Add0	2	200	0	3
S4	Periodic	cpu1	Add0	3	100	0	2
S5	Periodic	cpu1	Add0	2	200	0	3
S6	Periodic	cpu1	Add0	2	100	0	2
S7	Periodic	cpu1	Add0	2	1000	0	4

Figure 8: Tasks for seqgen2.c

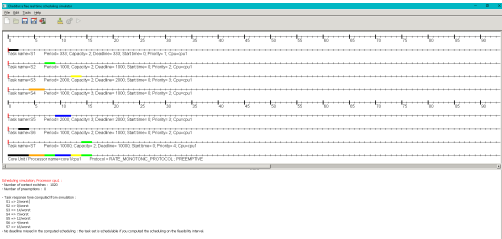


Figure 9: RM Execution for seqgen.c

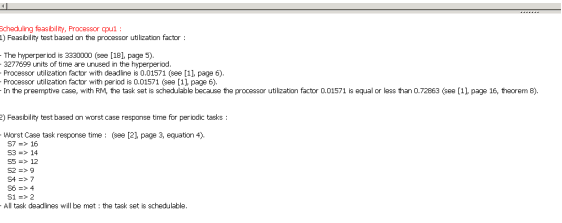


Figure 10: Feasibility test for seqgen.cs

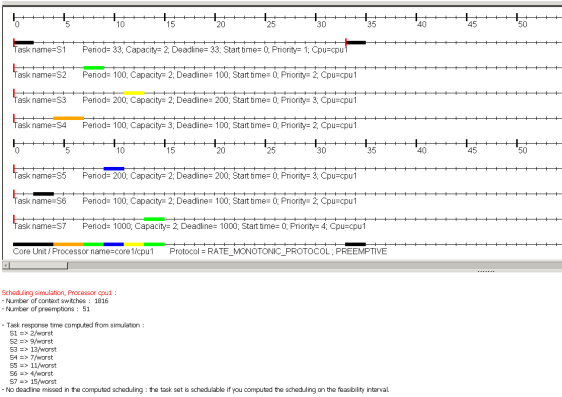


Figure 11: RM Execution for seqgen2.c

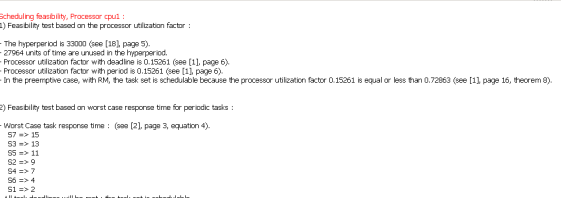


Figure 12: Feasibility test for seqgen2.c

2 Problem 2

2.1 Problem 2

Revise seqgen.c and seqgen2x.c to run under FreeRTOS on the DE1-SoC or TIVA board, by making each service a FreeRTOS task. Use the associated startup file in place of the existing

startup file in FreeRTOS. Use an ISR driven by the PIT hardware timer to release each task at the given rate (you could even put the sequencer in the ISR). Build and execute the new code. Determine the worst case execution time (WCET) for each service by printing or logging timestamps between two points in your code or by use of a profiling tool. Determine D, T, and C for each service and create an RM schedule in Cheddar using your WCET estimates. Calculate the Compare this with the results you achieved under Linux in (1).

2.2 Explanation and execution of code on target

The seqgen and the seqgen2x code was ported on the platform TIVA C Servies Launchpad with FreeRTOS as the operating system.

1. The 7 Services were created using the the function xTaskCreate
2. The sequencer was invoked in the ISR using the Timer 0 Timer A
3. UART was used to output on the serial terminal
4. The Tasks were assigned priorities using the RM policy in which the service with the highest frequency was assigned the highest and subsequent priorities in decreasing order of the frequency
5. The Sequencer and the services shared a semaphore among them for signalling
6. The timestamps was recorded using the API xTaskGetTickCount of the FreeRTOS kernel which updates its tick at the frequency configured in FreeRTOSConfig.h file
7. The printing on the serial terminal was protected by a common semaphore shared throughout the system
8. To handling the signal from the sequencer to services, the APIs xSemaphoreGive-FromISR and xSemaphoreTake was used
9. The Timer 0 was configured to occur periodically at the rate of the SystemClock divided the the Sequencer frequency which in this case was 30, 60 and 3000Hz.

2.3 Timestamps and WCET estimates

```
Worst Case Execution Time Service 1= 1
Worst Case Execution Time Service 2= 13
Worst Case Execution Time Service 3= 4
Worst Case Execution Time Service 4= 7
Worst Case Execution Time Service 5= 9
Worst Case Execution Time Service 6= 20
Worst Case Execution Time Service 7= 3
```

Figure 13: Output on serial terminal for seqgen

```
Worst Case Execution Time Service 1= 5
Worst Case Execution Time Service 2= 13
Worst Case Execution Time Service 3= 4
Worst Case Execution Time Service 4= 7
Worst Case Execution Time Service 5= 15
Worst Case Execution Time Service 6= 20
Worst Case Execution Time Service 7= 3
```

Figure 14: Output on serial terminal for seqgen

2.4 D,C and T Table for all services

Service	Frequency	Time Period T in ms	Deadline D, D=T	Priority	WCET	WCET in ms	Execution Time C	C/T
S1	30	33.33333333	33.33333333	1	1395938	1.395938	1.395938	0.041878
S2	10	100	100	2	1943073	1.943073	1.943073	0.019431
S4	10	100	100	2	2298021	2.298021	2.298021	0.02298
S6	10	100	100	2	1749011	1.749011	1.749011	0.01749
S3	5	200	200	3	1430416	1.430416	1.430416	0.007152
S5	5	200	200	3	1684479	1.684479	1.684479	0.008422
S7	1	1000	1000	4	2037500	2.0375	2.0375	0.002038
								0.119391

Figure 15: D,C and T Table for all services for Seqgen

Service	Frequency	Time Period T in ms	Deadline D, D=T	Priority	WCET in ms	Execution Time C	C/T
S1	30	33.33333333	33.33333333	1	5	5	0.15
S2	10	100	100	2	13	13	0.13
S4	10	100	100	2	4	4	0.04
S6	10	100	100	2	7	7	0.07
S3	5	200	200	3	15	15	0.075
S5	5	200	200	3	20	20	0.1
S7	1	1000	1000	4	3	3	0.003
							0.568

Figure 16: D,C and T Table for all services for Seqgen2X

2.5 Cheddar Schedule

```
Scheduling simulation, Processor TIVAFreeRTOS :
- Number of context switches : 11
- Number of preemptions : 0

- Task response time computed from simulation :
  Service1 => 2/worst
  Service2 => 55/worst
  Service3 => 35/worst
  Service4 => 32/worst
  Service5 => 77/worst
  Service6 => 66/worst
  Service7 => 80/worst

- No deadline missed in the computed scheduling : the task set seems to be schedulable.
```

Figure 17: RM Execution for seqgen.c

```
Scheduling simulation, Processor TIVAFreeRTOS :
- Number of context switches : 15
- Number of preemptions : 2

- Task response time computed from simulation :
  Service1 => 5/worst
  Service2 => 29/worst
  Service3 => 74/worst
  Service4 => 16/worst
  Service5 => 54/worst
  Service6 => 12/worst
  Service7 => 77/worst

- No deadline missed in the computed scheduling : the task set seems to be schedulable.
```

Figure 18: RM Execution for seqgen.c

2.6 % CPU Utilization

```
Scheduling feasibility, Processor TIVAFreeRTOS :
1) Feasibility test based on the processor utilization factor :
- The base period is 33333 ms (see [1], page 3).
- 33333 ms of time are needed in the base period.
- Processor utilization factor with deadline is 0.07031 (see [1], page 4).
- In the pessimistic case, with 80, the task set is schedulable because the processor utilization factor 0.07031 is equal or less than 0.72863 (see [1], page 36, theorem 6).

2) Feasibility test based on worst case task response time :
- Based on task response time : (see [2], page 3, equation 4).
  Service1 => 77
  Service2 => 55
  Service3 => 35
  Service4 => 32
  Service5 => 77
  Service6 => 66
  Service7 => 80
- All task deadlines will be met : the task set is schedulable.
```

Figure 19: Utilization of CPU for seqgen.c

```
Scheduling feasibility, Processor TIVAFreeRTOS :
1) Feasibility test based on the processor utilization factor :
- The base period is 33333 ms (see [1], page 3).
- 33333 ms of time are needed in the base period.
- Processor utilization factor with deadline is 0.06952 (see [1], page 4).
- In the pessimistic case, with 80, the task set is schedulable because the processor utilization factor 0.06952 is equal or less than 0.72863 (see [1], page 36, theorem 6).

2) Feasibility test based on worst case task response time :
- Based on task response time : (see [2], page 3, equation 4).
  Service1 => 77
  Service2 => 55
  Service3 => 35
  Service4 => 32
  Service5 => 77
  Service6 => 66
  Service7 => 80
- All task deadlines will be met : the task set is schedulable.
```

Figure 20: Utilization of CPU for seqgen2X.c

2.7 Comparision between Linux and FreeRTOS

The Utilization of CPU was found to be different on the Linux and FreeRTOS system. This might be due to the face that a synthetic load was introduced in FreeRTOS due to the resolution of the clocks. The number of context switches might also add to more utilization of CPU in FreeRTOS as the sequencer was handled in the interrupt

3 Problem 3

3.1 Problem

Revise seqgen.c from both previous systems to increase the sequencer frequency and all service frequencies by a factor of 100 (3000 Hz). Build and execute the code under Linux and FreeRTOS on your target boards as before. For both operating systems determine the worst case execution time (WCET) for each service by printing or logging timestamps between two points in your code or by use of a profiling tool. Determine D, T, and C for each service and create an RM schedule in Cheddar using your WCET estimates. Calculate the percentage

CPU utilization for these systems. Compare results between Linux and FreeRTOS in this higher-speed case.

3.2 Answer

3.2.1 Explanation and Execution of code

Linux seqgen.c at 3000Hz As described in section 1, the seqgen.c code runs the sequencer as a primary thread with all services running as parallel threads from it. The difference now is that the sequencer is being run at 3000Hz. This changes the services frequencies to the following: s1 = 300Hz, s2 = 100Hz, s3 = 50Hz, s4 = 100Hz, s5 = 50Hz, s6 = 100Hz and s7 = 10Hz. The figure below shows the execution of seqgen.c at 3000Hz.

```
patrick@rtnano:~/exercise5/e5$ sudo ./seqgen
Starting Sequencer Demo
System has 4 processors configured and 4 available.
Using CPUS=1 from total available.
Pthread Policy is SCHED_FIFO
PTHREAD SCOPE SYSTEM
rt_max_prio=99
rt_min_prio=1
Service threads will run on 1 CPU cores
pthread_create successful for service 1
pthread_create successful for service 2
Frame Sampler thread @ sec=0, msec=424
pthread_create successful for service 3
Time-stamp with Image Analysis thread @ sec=0, msec=424
Time-stamp with Image Analysis thread @ sec=0, msec=424
pthread_create successful for service 4
pthread_create successful for service 5
Time-stamp with Image Analysis thread @ sec=0, msec=424
pthread_create successful for service 6
Time-stamp with Image Analysis thread @ sec=0, msec=424
pthread_create successful for service 7
Start sequencer
pthread_create successful for sequencer service 0
Sequencer thread @ sec=0, msec=425
Time-stamp with Image Analysis thread @ sec=0, msec=424
Time-stamp with Image Analysis thread @ sec=0, msec=424
Worst Case Execution Time Service = 1 = 299166
Worst Case Execution Time Service = 4 = 443385
Worst Case Execution Time Service = 2 = 400261
Worst Case Execution Time Service = 3 = 258593
Worst Case Execution Time Service = 6 = 619010
Worst Case Execution Time Service = 7 = 685313
Worst Case Execution Time Service = 5 = 705052
TEST COMPLETE
```

Figure 21: seqgen.c at 3kHz

Cheddar and CPU Utilization As in section 1, the values of C for cheddar simulations were rounded up. This change allows the use of integers for Cheddar simulation and guarantees that any c value lower than those used will be a viable solution. The CPU utilization for the actual WCET values was 28.649% and the Utilization with the rounded values was 71%. These calculations and the Cheddar simulation results are pictured below.

3.3 D,C and T Table for all services for Free RTOS

Service	Frequency	Period T	Deadline D	Priority	WCET	WCET ms	Execution Time C	CPU Utilization WCET	CU Utilization Rounded Values
S1	300	3.333333333	3.333333333	1	299166	0.299166	1	0.0897498	0.3
S2	100	10	10	2	443385	0.443385	1	0.0443385	0.1
S3	50	20	20	3	400261	0.400261	1	0.02001305	0.05
S4	100	10	10	2	258593	0.258593	1	0.0258593	0.1
S5	50	20	20	3	619010	0.61901	1	0.0309505	0.05
S6	100	10	10	2	685313	0.685313	1	0.0685313	0.1
S7	10	100	100	4	705052	0.705052	1	0.0705052	0.1
							Total CPU Utilization:	28.649%	71.000%

Figure 22: Utilization Calculations at 3kHz

3.4 Cheddar schedule and % CPU Utilization Linux

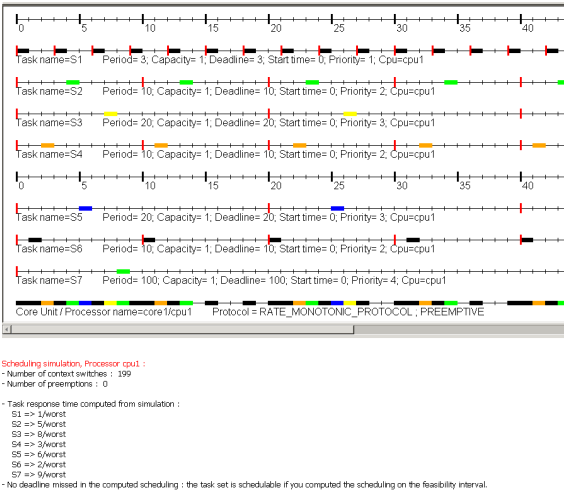


Figure 23: Cheddar Simulation 3kHz

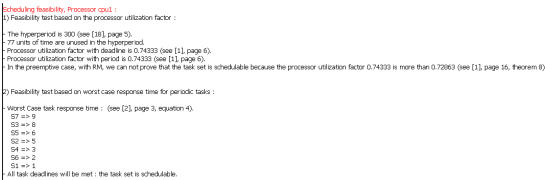


Figure 24: Cheddar Feasibility 3kHz

3.5 D,C and T Table for all services for Free RTOS

Service	Frequency	Time Period T in ms	Deadline D, D _{min}	Priority	WCET in ms	Execution C/T	
S1	300	3.333333333	3.333333333	1	2	2	0.6
S2	100	10	10	2	20	20	2
S4	100	10	10	2	29	29	2.9
S6	100	10	10	2	10	10	1
S3	50	20	20	3	2	2	0.1
S5	50	20	20	3	11	11	0.55
S7	10	100	100	4	3	3	0.03
							7.18

Figure 25: D,T and C Values

3.6 Cheddar schedule and % CPU Utilization

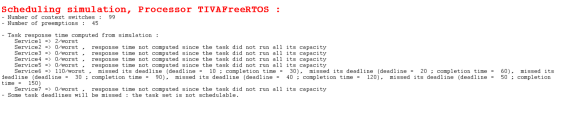


Figure 26: Free RTOS Cheddar Simulation 3kHz

Figure 27: Free RTOS Cheddar Feasibility 3kHz

3.7 Comparision between Linux and Free RTOS

The services are not schedule in FreeRTOS due to the processing speed. This might also be due to the number of context switches

4 Appendix- Code

Code Repository

 [Github Link](#)